Intro
○○○○○

Examples
○○○○

Propositional logic
○○○○○○○○○○○○○

Classic
○○

Cut
○○

Quantifiers
○○○

Equality
○○○○○○○○○○○○

Induction
○○○

Poor Man's Axi: DPL-like proofs in Type Theory

Wojciech Kołowski

Layers in Poor Man's Axi

In our original proposal of Poor Man's Axi, the language was split into two layers:

- A programming layer which consists of a strongly-typed functional programming language based on the Simply Typed Lambda Calculus.

- A logical layer which consists of second-order classical logic with equality and some primitives for reasoning by cases and induction (note that it's not full second-order logic, because it only allows quantification over functions, but not over propositions, predicates or relations).

Intro
○●○○○

Examples
○○○○

Propositional logic
○○○○○○○○○○○○

Classic
○○

Cut
○○

Quantifiers
○○○

Equality
○○○○○○○○○○○○

Induction
○○○

Logic without proofterms

The logic was presented with a bunch of judgements, the most important being the true proposition judgement $\Gamma \mid \Delta \vdash P$. While this presentation does a good job of explaining what the logic is like, it does not address the problem of writing proofs from the perspective of the programmer.

Logic with proofterms

To deal with this problem, we introduce proofterms (also known as proof certificates), which are expressions that will serve as the "proofs" that the user will write. The manipulation of judgements which establish their correctness (which we will describe soon) is left to the language's proof checker.

Proofterms

Proofterms (here P are propositions, t are terms, x are variables):

$e ::=$

$P \mid$ **true** $\mid$ **exfalso** e

assume P **in** $e \mid$ **modus-ponens** e_1 $e_2 \mid$

suppose-absurd P **in** $e \mid$ **absurd** e_1 $e_2 \mid$

both e_1 $e_2 \mid$ **left-and** $e \mid$ **right-and** $e \mid$

equivalence e_1 $e_2 \mid$ **left-iff** $e \mid$ **right-iff** $e \mid$

left-either P $e \mid$ **right-either** P $e \mid$

constructive-dilemma e_1 e_2 $e_3 \mid$

double-negation $e \mid$

$e_1; e_2 \mid$

pick-any x **in** $e \mid$ **specialize** e **with** $t \mid$

exists t **such that** $e \mid$ **pick-witness** x **for** e_1 **in** $e_2 \mid$

refl $t \mid$ **rewrite** e_1 **in** $e_2 \mid$

case t **of** $(\text{inl } a \rightarrow e_1, \text{inr } b \rightarrow e_2) \mid$

Proofterms – overview

A quick glance at the above grammar:

- First, propositions are included in the syntactic category of proofterms, because we will refer to an assumption P with P itself.

- Then we have some proofterms for dealing with propositional logic, divided as usual into introduction and elimination forms.

- **double-negation** is our way of making the logic classical.

- $e_1; e_2$ is a cut (which, in DPLs, corresponds to proof composition).

- Then we have some proofterms to deal with quantifiers.

- **refl** and **rewrite** are primitives for dealing with equality.

- Finally, **case** will be used for reasoning by cases on terms of type $A + B$.

Examples

Before we jump right into the rules, let's see some examples. The point of the first two is to show that our proofterms look very similar to proofs from DPLs (in particular, to proofs in $\mathcal{NDL}$ presented in chapters 4 and 6 of the DPL thesis). The third example shows how simple reasoning about programs might look like.

Example – propositional logic

Theorem: $(P \Rightarrow Q) \Rightarrow (Q \Rightarrow R) \Rightarrow P \Rightarrow R$.

Proof:
assume $P \Rightarrow Q$ **in**
 assume $Q \Rightarrow R$ **in**
  assume P **in**
   modus-ponens $(P \Rightarrow Q)$ P;
   modus-ponens $(Q \Rightarrow R)$ Q

The proof looks the same as the $\mathcal{NDL}$ one (DPL thesis, chapter 4, page 71), except that we don't have **begin** and **end**.

Example – first-order logic

Theorem: $(\forall x : A.\ P\ x \wedge Q\ x) \Rightarrow (\forall x : A.\ P\ x) \wedge (\forall x : A.\ Q\ x)$

Proof:
assume $\forall x : A.\ P\ x \wedge Q\ x$ **in**
  pick-any y **in**
    specialize $\forall x : A.\ P\ x \wedge Q\ x$ **with** y;
    left-and $P\ y \wedge Q\ y$;
  pick-any y **in**
    specialize $\forall x : A.\ P\ x \wedge Q\ x$ **with** y;
    right-and $P\ y \wedge Q\ y$;
  both $(\forall y : A.\ P\ y)\ (\forall y : A.\ Q\ y)$

Again, the proof looks the same as the $\mathcal{NDL}$ one (DPL thesis, chapter 6, page 156), except we use indentation instead of **begin** and **end**.

Example – proof about a program

Consider the following program:
swap : $A + B \rightarrow B + A$:= $\lambda x.$case x of $(\lambda a.$inr $a, \lambda b.$inl $b)$

It's pretty clear that this function is involutive.
Theorem: $\forall x : A + B.$ swap (swap x) = x

Proof:
pick-any x **in**
  case x **of** (inl $a \rightarrow$ **refl** a, inr $b \rightarrow$ **refl** b)

The proof has the same structure as the proofterm you would write
in Coq, except for the syntactic differences.

Correct proof judgement

We now proceed to describe proofterms in much greater detail. We will modify the language in the following way: we throw away the true proposition judgement $\Gamma \mid \Delta \vdash P$ and replace it with the correct proof judgement $\Gamma \mid \Delta \vdash e : P$, which should be read: in typing context Γ and assumption context Δ, e is a proof of P.

The other judgements we need in our logic, the valid assumption context judgement $\Gamma \vdash \Delta$ valid and the well-formed proposition judgement $\Gamma \vdash P$ prop, are unchanged. Note that we will set up the system so that $\Gamma \mid \Delta \vdash e : P$ entails $\Gamma \vdash \Delta$ valid and $\Gamma \vdash P$ prop as sanity checks.

Intro
○○○○○

Examples
○○○○

Propositional logic
○●○○○○○○○○○○○

Classic
○○

Cut
○○

Quantifiers
○○○

Equality
○○○○○○○○○○○○

Induction
○○○

Assumptions

$$\frac{\Gamma \;\vdash \Delta \; \texttt{valid} \quad P \in \Delta}{\Gamma \mid \Delta \vdash P : P} \text{Ass}$$

The basic rule of our logic is that we can use assumptions from the assumption context. Note that here P denotes both the proposition and its proof.

True and False

$$\frac{\Gamma \;\vdash \Delta \;\texttt{valid}}{\Gamma \mid \Delta \vdash \textbf{true} : \top}\text{True-Intro}$$

$$\frac{\Gamma \mid \Delta \vdash e : \bot}{\Gamma \mid \Delta \vdash \textbf{exfalso}\; e : P}\text{False-Elim}$$

true proves the true proposition $\top$.

exfalso e proves any proposition P whatsoever, provided that e proves the false proposition $\bot$.

Implication

$$\frac{\Gamma \mid \Delta, P \vdash e : Q}{\Gamma \mid \Delta \vdash \textbf{assume } P \textbf{ in } e : P \Rightarrow Q} \text{Impl-Intro}$$

$$\frac{\Gamma \mid \Delta \vdash e_1 : P \Rightarrow Q \quad \Gamma \mid \Delta \vdash e_2 : P}{\Gamma \mid \Delta \vdash \textbf{modus-ponens } e_1 \ e_2 : Q} \text{Impl-Elim}$$

assume P **in** e proves $P \Rightarrow Q$, provided that e proves Q in the assumption context extended with P.

modus-ponens $e_1 \ e_2$ proves Q, provided that e_1 proves $P \Rightarrow Q$ and e_2 proves P.

Negation

$$\frac{\Gamma \mid \Delta, P \vdash e : \bot}{\Gamma \mid \Delta \vdash \textbf{suppose-absurd } P \textbf{ in } e : \neg P}\text{Not-Intro}$$

$$\frac{\Gamma \mid \Delta \vdash e_1 : \neg P \quad \Gamma \mid \Delta \vdash e_2 : P}{\Gamma \mid \Delta \vdash \textbf{absurd } e_1 \ e_2 : \bot}\text{Not-Elim}$$

suppose-absurd P **in** e proves $\neg P$ provided that e proves $\bot$ in the context extended with P as an assumption.

absurd $e_1 \ e_2$ proves $\bot$ provided that e_1 proves $\neg P$ and e_2 proves P.

Negation – discussion

Note that there are some differences from $\mathcal{NDL}$ (DPL thesis, chapter 4, page 65). First, **suppose-absurd** is a primitive, instead of being syntax sugar. Second, the order of arguments of **absurd** is flipped. The reason for this is that we follow a presentation of logic derived from the Curry-Howard correspondence.

However, what we do is not *the* usual Curry-Howard presentation, because in our system negation is primitive. The usual way to deal with negation would be to *define* $\neg P$ as an abbreviation for $P \Rightarrow \bot$, so that we would deal with negation by using **assume** instead of **suppose-absurd** and **modus-ponens** instead of **absurd**.

Intro
00000

Examples
0000

Propositional logic
000000●00000

Classic
00

Cut
00

Quantifiers
000

Equality
00000000000

Induction
000

Conjunction

$$\frac{\Gamma \mid \Delta \vdash e_1 : P \quad \Gamma \mid \Delta \vdash e_2 : Q}{\Gamma \mid \Delta \vdash \textbf{both } e_1 \ e_2 : P \wedge Q} \text{And-Intro}$$

$$\frac{\Gamma \mid \Delta \vdash e : P \wedge Q}{\Gamma \mid \Delta \vdash \textbf{left-and } e : P} \text{And-Elim-L}$$

$$\frac{\Gamma \mid \Delta \vdash e : P \wedge Q}{\Gamma \mid \Delta \vdash \textbf{right-and } e : Q} \text{And-Elim-R}$$

both e_1 e_2 proves $P \wedge Q$ provided that e_1 proves P and e_2 proves Q. **left-and** e proves P provided that e proves $P \wedge Q$. **right-and** e proves Q provided that e proves $P \wedge Q$.

Biconditional

$$\frac{\Gamma \mid \Delta \vdash e_1 : P \Rightarrow Q \quad \Gamma \mid \Delta \vdash e_2 : Q \Rightarrow P}{\Gamma \mid \Delta \vdash \textbf{equivalence } e_1 \; e_2 : P \Leftrightarrow Q} \text{Iff-Intro}$$

$$\frac{\Gamma \mid \Delta \vdash e : P \Leftrightarrow Q}{\Gamma \mid \Delta \vdash \textbf{left-iff } e : P \Rightarrow Q} \text{Iff-Elim-L}$$

$$\frac{\Gamma \mid \Delta \vdash e : P \Leftrightarrow Q}{\Gamma \mid \Delta \vdash \textbf{right-iff } e : Q \Rightarrow P} \text{Iff-Elim-R}$$

equivalence $e_1 \; e_2$ proves $P \Leftrightarrow Q$ provided that e_1 proves $P \Rightarrow Q$ and e_2 proves $Q \Rightarrow P$. **left-iff** e proves $P \Rightarrow Q$ provided that e proves $P \Leftrightarrow Q$. **right-iff** e proves $Q \Rightarrow P$ provided that e proves $P \Leftrightarrow Q$.

Biconditional – discussion

Our presentation of the biconditional is exactly the same as in $\mathcal{NDL}$ (DPL thesis, chapter 4, page 65), although it differs from the usual Curry-Howard presentation: in our system $P \Leftrightarrow Q$ is a primitive, whereas in the usual Curry-Howard presentation it would be defined as $(P \Rightarrow Q) \wedge (Q \Rightarrow P)$, so that we would use **both** instead of **equivalence**, **left-and** instead of **left-iff** and **right-and** instead of **right-iff**.

Disjunction

$$\frac{\Gamma \mid \Delta \vdash e : P}{\Gamma \mid \Delta \vdash \textbf{left-either } Q \ e : P \vee Q} \text{Or-Intro-L}$$

$$\frac{\Gamma \mid \Delta \vdash e : Q}{\Gamma \mid \Delta \vdash \textbf{right-either } P \ e : P \vee Q} \text{Or-Intro-R}$$

$$\frac{\Gamma \mid \Delta \vdash e_1 : P \vee Q \quad \Gamma \mid \Delta \vdash e_2 : P \Rightarrow R \quad \Gamma \mid \Delta \vdash e_3 : Q \Rightarrow R}{\Gamma \mid \Delta \vdash \textbf{constructive-dilemma } e_1 \ e_2 \ e_3 : R} \text{Or-Elim}$$

left-either Q e proves $P \vee Q$ provided that e proves P.
right-either P e proves $P \vee Q$ provided that e proves Q.
constructive-dilemma e_1 e_2 e_3 proves R provided that e_1 proves
$P \vee Q$, e_2 proves $P \Rightarrow R$ and e_3 proves $Q \Rightarrow R$.

Disjunction – annotations

The first argument of **left-either** Q e, i.e. Q, is a proposition, not a proofterm, and serves as an annotation that clarifies what its conclusion is. **right-either** P e is analogous.

This choice makes our system very close to $\mathcal{NDL}$ (DPL thesis, chapter 4, page 65), and also to the "intrinsic" presentation of the Simply Typed Lambda Calculus, but it is not the only one: we could do without the annotation, which would make the system less similar to $\mathcal{NDL}$ and more to the "extrinsic" presentation of Simply Typed Lambda Calculus.

Intro
00000
Examples
0000
Propositional logic
0000000000●0
Classic
00
Cut
00
Quantifiers
000
Equality
0000000000
Induction
000

Disjunction – choices

We also have a choice regarding the second and third arguments of
constructive-dilemma e_1 e_2 e_3. To keep close to $\mathcal{NDL}$, they
both must prove implications ($P \Rightarrow R$ for e_2 and $Q \Rightarrow R$ for e_3),
but perhaps it would be more comfortable for the user if they only
needed to prove the conclusion R in a context extended with the
appropriate assumption, so that he would be spared from having to
write **assume** twice (once in e_2 and once in e_3). The alternative
rule is as follows:

$$\frac{\Gamma \mid \Delta \vdash e_1 : P \vee Q \quad \Gamma \mid \Delta, P \vdash e_2 : R \quad \Gamma \mid \Delta, Q \vdash e_3 : R}{\Gamma \mid \Delta \vdash \textbf{constructive-dilemma } e_1 \ e_2 \ e_3 : R}\text{OR-ELIM-ALT}$$

Classic logic

Up until now, our logic is entirely constructive, closely following
the Curry-Howard correspondence, although it differed slightly with
regards to the treatment of negation and biconditional.

To make the logic classical, we need to extend it with one of the
laws of classical logic that is strong enough to entail all the others.
In our previous presentation of Axi we chose the Law of Excluded
Middle, but this time, to keep close to $\mathcal{NDL}$, we will use Double
Negation Elimination.

In a richer language and logic, the way to go would probably be to
assert some version of the Axiom of Choice.

Double Negation Elimination

$$\frac{\Gamma \mid \Delta \vdash e : \neg\neg P}{\Gamma \mid \Delta \vdash \textbf{double-negation } e : P}\textsc{Classic}$$

double-negation e proves P provided that e proves $\neg\neg P$.

Backward and forward proofs

The proofterms we introduced so far only allowed for comfortable "backward" proofs, i.e. ones where we go from the conclusion to the assumptions. The other style, "forward" proofs, are also possible, but doing them is not comfortable.

To amend this problem, we introduce a new rule, called the CUT rule. In DPLs, it corresponds to proof composition (and in programming languages, to let-expressions, although there is no variable binding).

The CUT rule increases the comfort of doing forward proofs in our system, but does not add any expressive power, because $e_1; e_2$ can be desugared to **modus-ponens** (**assume** P **in** e_2) e_1.

Cut rule

$$\frac{\Gamma \mid \Delta \vdash e_1 : P \quad \Gamma \mid \Delta, P \vdash e_2 : Q}{\Gamma \mid \Delta \vdash e_1; e_2 : Q} \text{Cut}$$

$e_1; e_2$ proves Q provided that e_1 proves P and e_2 proves Q in the context extended with the assumption P.

Quantifiers

Our syntax for quantifiers closely follows $\mathcal{NDL}$ (DPL thesis, chapter 6, page 150), except for the introduction rule of the existential quantifier, in which the argument order is flipped (and which, for this reason, uses different keywords).

Our rules, however, are subtly different in that we allow arbitrary proofs as arguments to **specialize**, **exists** and **pick-witness**, whereas $\mathcal{NDL}$ only allows formulas to appear as arguments in these places. Therefore, our approach is a slight generalization, which also closely aligns in spirit with the Curry-Howard correspondence.

Universal quantifier

$$\frac{\Gamma, y : A \mid \Delta \vdash e : P\,[x := y]}{\Gamma \mid \Delta \vdash \textbf{pick-any } y \textbf{ in } e : \forall x : A.\, P}\text{Forall-Intro}$$

$$\frac{\Gamma \mid \Delta \vdash e : \forall x : A.\, P \quad \Gamma \vdash t : A}{\Gamma \mid \Delta \vdash \textbf{specialize } e \textbf{ with } t : P\,[x := t]}\text{Forall-Elim}$$

pick-any y **in** e proves $\forall x : A.\, P$ provided that e proves P in which y was substituted for x.

specialize e **with** t proves P in which t was substituted for x, provided that e proves $\forall x : A.\, P$ and t is a term of type A.

Existential quantifier

$$\frac{\Gamma \vdash t : A \quad \Gamma \mid \Delta \vdash e : P\,[x := t]}{\Gamma \mid \Delta \vdash \textbf{exists } t \textbf{ such that } e : \exists x : A.\, P}\text{Exists-Intro}$$

$$\frac{\Gamma \vdash R \text{ prop} \quad \Gamma \mid \Delta \vdash e_1 : \exists x : A.\, P \quad \Gamma, y : A \mid \Delta, P\,[x := y] \vdash e_2 : R}{\Gamma \mid \Delta \vdash \textbf{pick-witness } y \textbf{ for } e_1 \textbf{ in } e_2 : R}\text{Exis}$$

exists t **such that** e proves $\exists x : A.\, P$ provided that t is a term of
type A and e proves P in which t was substituted for x.

pick-witness y **for** e_1 **in** e_2 proves R, provided that e_1 proves
$\exists x : A.\, P$ and e_2 proves R in the typing context extended with
$y : A$ and the assumption context extended with P (in which y was
substituted for x).

Computational equality – refresher

The computational equality judgement $\Gamma \vdash t_1 \equiv t_2 : A$ (sometimes also called judgemental equality, or simply convertibility) is the way in which we express computation in our language. Its intuitive meaning is that t_1 and t_2 compute to the same result (although the rules we gave in the previous slides don't look like this at all).

Computational equality is a part of the programming layer of our language and does not depend on the logical layer in any way. In particular, it does not have to change to accomodate proofterms. Computational equality is decidable and its use in proofs is handled entirely by the proof checker.

Propositional equality – refresher

Propositional equality $t_1 =_A t_2$ is not a judgement, but, as the name suggests, a proposition. It belongs to the logical layer of our language and the introduction of proofterms does affect it – we will need to change our rules to accomodate them. It is not decidable and all proofs involving propositional equality are up to the programmer.

Equality – overview

Our treatment of equality is fundamentally type-theoretic, with the conversion rule, which connects computational equality to propositional equality, playing the central role. The end product, however, is very similar to $\mathcal{NDL}$ (DPL thesis, chapter 6, page 151).

We also include a rule that asserts function extensionality, i.e. the principle that two functions are equal when they are equal for all arguments. This rule is not present in $\mathcal{NDL}$ because its language of terms doesn't have function types.

Conversion rule

$$\frac{\Gamma, x : A \vdash P \text{ prop} \quad \Gamma \vdash t_1 \equiv t_2 : A \quad \Gamma \mid \Delta \vdash e : P\left[x := t_1\right]}{\Gamma \mid \Delta \vdash e : P\left[x := t_2\right]}\text{Conv}$$

If e proves P with t_1 substituted for x and t_1 is computationally equal to t_2 at type A, then e also proves P with t_2 substituted for x.

Conversion rule – discussion

The main point of the conversion rule is to facilitate proving equalities which hold by computation. In contrast to DPLs, which would require doing some manual labor here (or calling powerful methods for help), in our system proving such equalities is free, or even more than free – the programmer doesn't need to do anything to benefit from the conversion rule, not even use the conversion rule. This is because computational equality is decidable, so using the conversion rule at the right time is entirely up to the proof checker.

Equality introduction and elimination

$$\frac{\Gamma \;\vdash\; \Delta \;\texttt{valid} \quad \Gamma \vdash t : A}{\Gamma \mid \Delta \vdash \textbf{refl } t : t =_A t}\text{Eq-Intro}$$

$$\frac{\Gamma, x : A \vdash P \;\texttt{prop} \quad \Gamma \mid \Delta \vdash e : t_1 =_A t_2 \quad \Gamma \mid \Delta \vdash e' : P\left[x := t_1\right]}{\Gamma \mid \Delta \vdash \textbf{rewrite } e \textbf{ in } e' : P\left[x := t_2\right]}\text{Eq-Elim}$$

refl t proves that t is equal to itself, provided that it is well-typed at some type A.

rewrite e **in** e' proves P with t_2 substituted for x, provided that e proves $t_1 =_A t_2$ (for some type A) and that e' proves P with t_1 substituted for x.

Intro
○○○○○

Examples
○○○○

Propositional logic
○○○○○○○○○○○○○

Classic
○○

Cut
○○

Quantifiers
○○○

Equality
○○○○○○●○○○○

Induction
○○○

refl – discussion

refl t might look like a weakling, a puny little proofterm only able to prove $t = t$, but in reality it is an incantation that summons the invincible juggernaut – the conversion rule. If t is computationally equal to t', then, by the conversion rule, to prove $t = t'$ it suffices to prove $t = t$, which is precisely what **refl** t accomplishes.

rewrite – discussion

rewrite is the workhorse for equational proofs, corresponding to $\mathcal{NDL}$'s **leibniz** (DPL thesis, chapter 6, page 151), except that, as for quantifiers, we allow its arguments to be proofs, whereas $\mathcal{NDL}$ requires them to be formulas and terms whose equality is already in the assumption base.

Compared to Athena, using **rewrite** e **in** e' is similar to performing a single step of chaining, with e being the justifier and e' being the rest of the chain.

Compared to Coq, **rewrite** e **in** e' looks most similar to Coq's tactic rewrite $e; e'$, which rewrites an equation e and then goes on to e', which is the rest of the proof.

Function extensionality

$$\frac{\Gamma \mid \Delta \vdash e : \forall x : A.\ f\ x =_B g\ x}{\Gamma \mid \Delta \vdash \textbf{funext}\ e : f =_{A \to B} g}\text{Funext}$$

funext e proves $f = g$ provided that e proves $\forall x : A.\ f\ x = g\ x$.

Function extensionality – discussion

A separate rule for function extensionality is needed, because
extensionality, despite being a very basic, useful and desirable
principle, does not follow from anything else in the system: neither
from the intuitionistic core, nor the classical **double-negation**, nor
the rules for quantifiers, nor even the conversion rule, even though
function extensionality is present at the level of computational
equality. This is a well known phenomenon in type theory, with
which $\mathcal{NDL}$ and other DPLs don't have to deal because they
don't have function types.

Function extensionality – alternative rule

funext e is in some sense "higher-order", i.e. e is expected to be a proof of a universal statement, which means that usually it will be of the form **pick-any** x **in** e'. This suggests an alternative, "first-order" version of the rule for function extensionality, which has this variable binding baked-in:

$$\frac{\Gamma, x : A \mid \Delta \vdash e : f\ x =_B g\ x}{\Gamma \mid \Delta \vdash \textbf{funext}\ x\ \textbf{in}\ e : f =_{A \to B} g}\ \text{Funext-Alt}$$

Reasoning by cases (for sums)

$$\frac{\Gamma, x : A \vdash P \text{ prop} \quad \Gamma \vdash t : A + B \quad \begin{array}{l} \Gamma, a : A \mid \Delta \vdash e_1 : P[x := \text{inl } a] \\ \Gamma, b : B \mid \Delta \vdash e_2 : P[x := \text{inr } b] \end{array}}{\Gamma \mid \Delta \vdash \textbf{case } t \textbf{ of } (\text{inl } a \to e_1, \text{inr } b \to e_2) : P[x := t]}$$

case t **of** $(\text{inl } a \to e_1, \text{inr } b \to e_2)$ proves P with t substituted
for x, provided that t is of type $A + B$, that e_1 proves P with
$\text{inl } a$ substituted for x in a typing context extended with $a : A$,
and that e_2 proves P with $\text{inr } b$ substituted for x in a typing
context extended with $b : B$.

Reasoning by cases (for sums) – discussion

From the grammar and typing rules given in Poor Man's Axi slides,
we know that terms of type $A + B$ can be of only two forms: either
inl a for some $a : A$, or inr b for some $b : B$.

We know this, but our logic doesn't automagically know this, so
we need to tell it by adding an appropriate proofterm and a rule to
handle it. This is what **case** accomplishes.

Induction

Sums are not the only type for which we need to add a principle
for reasoning by cases – there are also products, the unit type, and
the empty type, after all.

But for more complicated types, like the natural numbers or lists,
reasoning by cases won't suffice, and so we would also need to add
proofterms and rules for proofs by induction. We won't do this
now for lack of space and time, but it's not much more difficult
than handling reasoning by cases for sums.